

# Dataflow under the von Neumann Machine: A New Paradigm for Computing Systems

**Xian-He Sun**

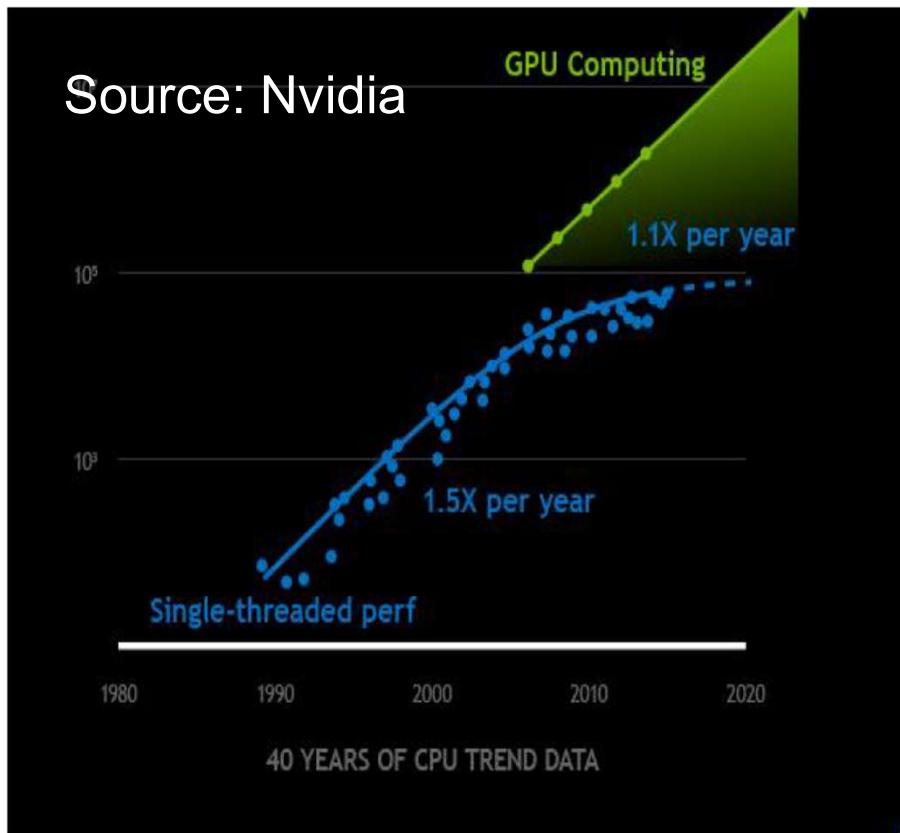
**Gnosis Research Center** for accelerating data-driven discovery

Illinois Institute of Technology  
[sun@iit.edu](mailto:sun@iit.edu)

**PERMAVOST2025**

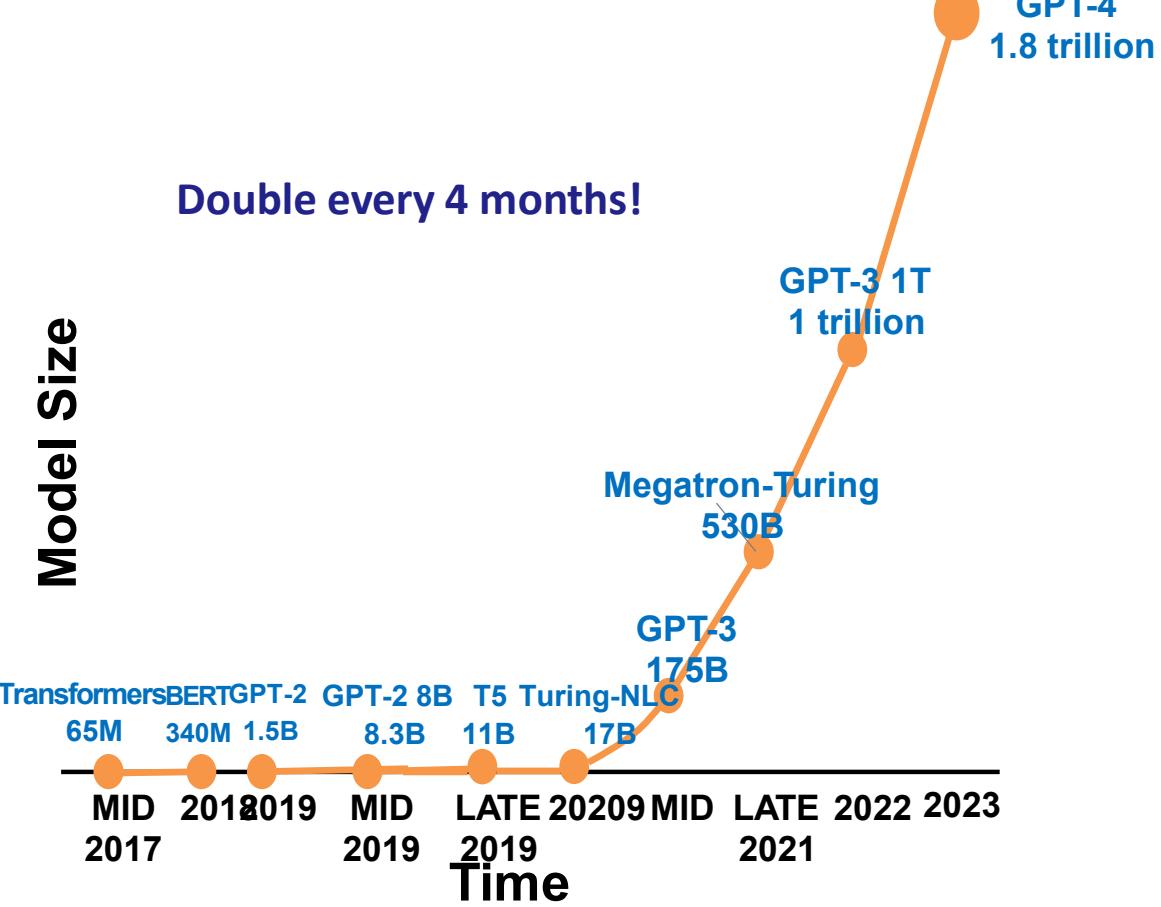


# The Computing & Data Disparity in AI



The new Moors' Law

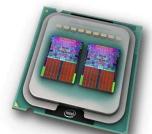
Courtesy of Tuo-Hung (Alex)Hou





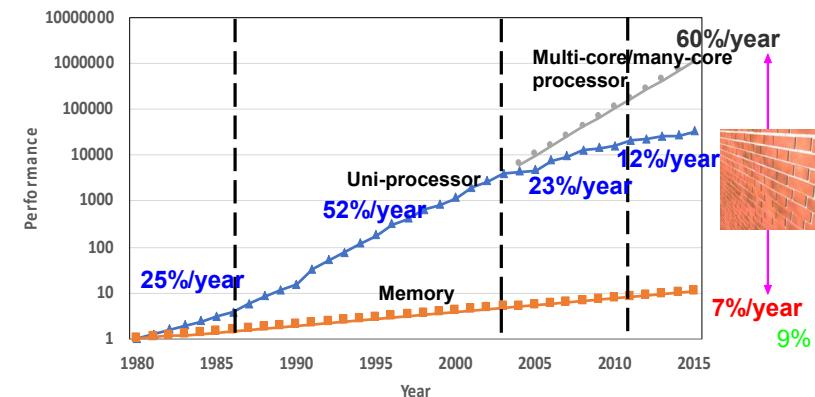
# Why Data Access ? The Memory-wall Problem

- Processor performance increases rapidly
  - Uni-processor: ~52% until 2004
  - Aggregate multi-core/many-core processor performance even higher since 2004
- Memory: ~7% per year
  - Storage: ~6% per year
- Processor-memory speed gap keeps increasing

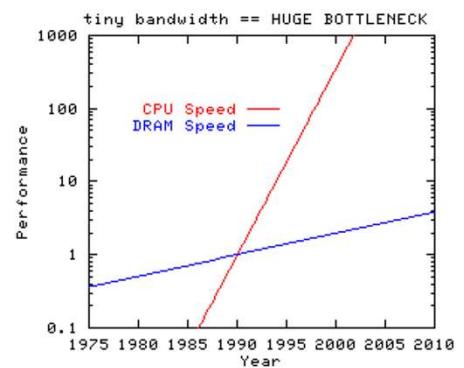


Source: Intel

Leiserson, Charles E., et al. "There's plenty of room at the Top: What will drive computer performance after Moore's law?." *Science* 368.6495 (2020).



Source: OCZ



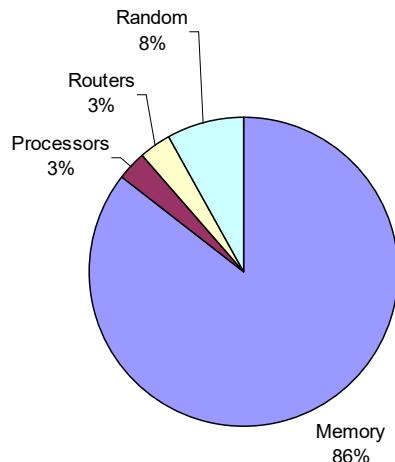
## Memory-bounded speedup (1990), Memory wall problem (1994)

X.-H. Sun, and L. Ni, "Scalable Problems and **Memory-Bounded Speedup**," Journal of Parallel and Distributed Computing, Vol. 19, 1993.  
X.-H. Sun, and L. Ni, "Another View of Parallel Speedup," Proc. of ACM-IEEE Supercomputing'90, NY, Nov. 1990

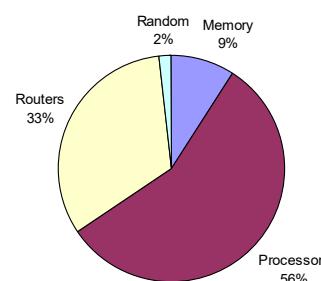


# Impact of Memory-Bound (-Wall)

Silicon Area Distribution



Power Distribution



- Modern microprocessors such as the Pentium Pro, Alpha 21164, Strong Arm SA110, and Longson-3A use 80% or more of their transistors for the on-chip cache

- 1989 the first Intel processor with on-chip L1 cache was Intel 486, 8KB size
- 1995 the first Intel processor with on-chip L2 cache was Intel Pentium Pro, 256KB size
- 2003 the first Intel processor with on-chip L3 cache was Intel Itanium 2, 6MB size
- 1980: no cache in microprocessor; 2010: 3-level cache on chip, 4-level cache off chip

Source: Computer Architecture A Quantitative Approach

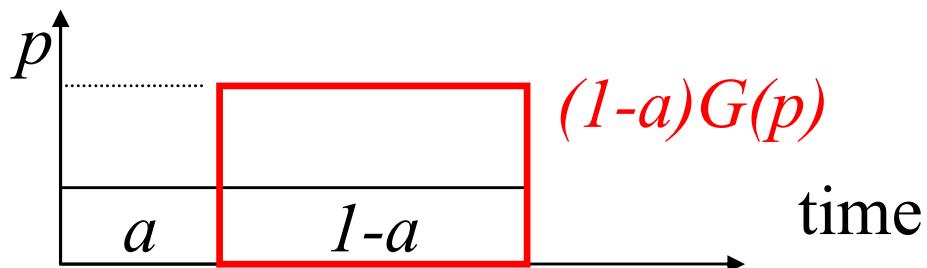
Courtesy of Peter Kogge, UND



# Memory-Bound & Scaling Function (Sun-Ni's Law)



Xian-He Sun



$$\text{Speedup}_{MB} = \frac{\text{Work}(p)/\text{Time}(p)}{\text{Work}(1)/\text{Time}(1)} = \frac{\alpha + (1-\alpha)G(p)}{\alpha + (1-\alpha)G(p)/p} \quad W^* = G(p * M)$$

## Scalability of Multicore

$$\frac{\frac{w_c}{\text{perf}(r)} + \frac{w_p}{m \cdot \text{perf}(r)}}{\frac{w_c}{\text{perf}(r)} + \frac{w_p}{\text{perf}(r)}} = \frac{\frac{w_c + m \cdot w_p}{w_c + w_p}}{(1 - f') + mf'} = \frac{w_p}{w_c + w_p}$$
$$f' = \frac{w_p}{w_c + w_p}$$

X.-H. Sun and Y. Chen, "Reevaluating Amdahl's Law in the Multicore Era," Journal of Parallel and Distributed Computing, vol. 70, no. 2, pp. 183-188, Feb. 2010

Based on Sun-Ni's law **Multicore is scalable**, if data access time is fixed and does not increase with the amount of work & the number of cores



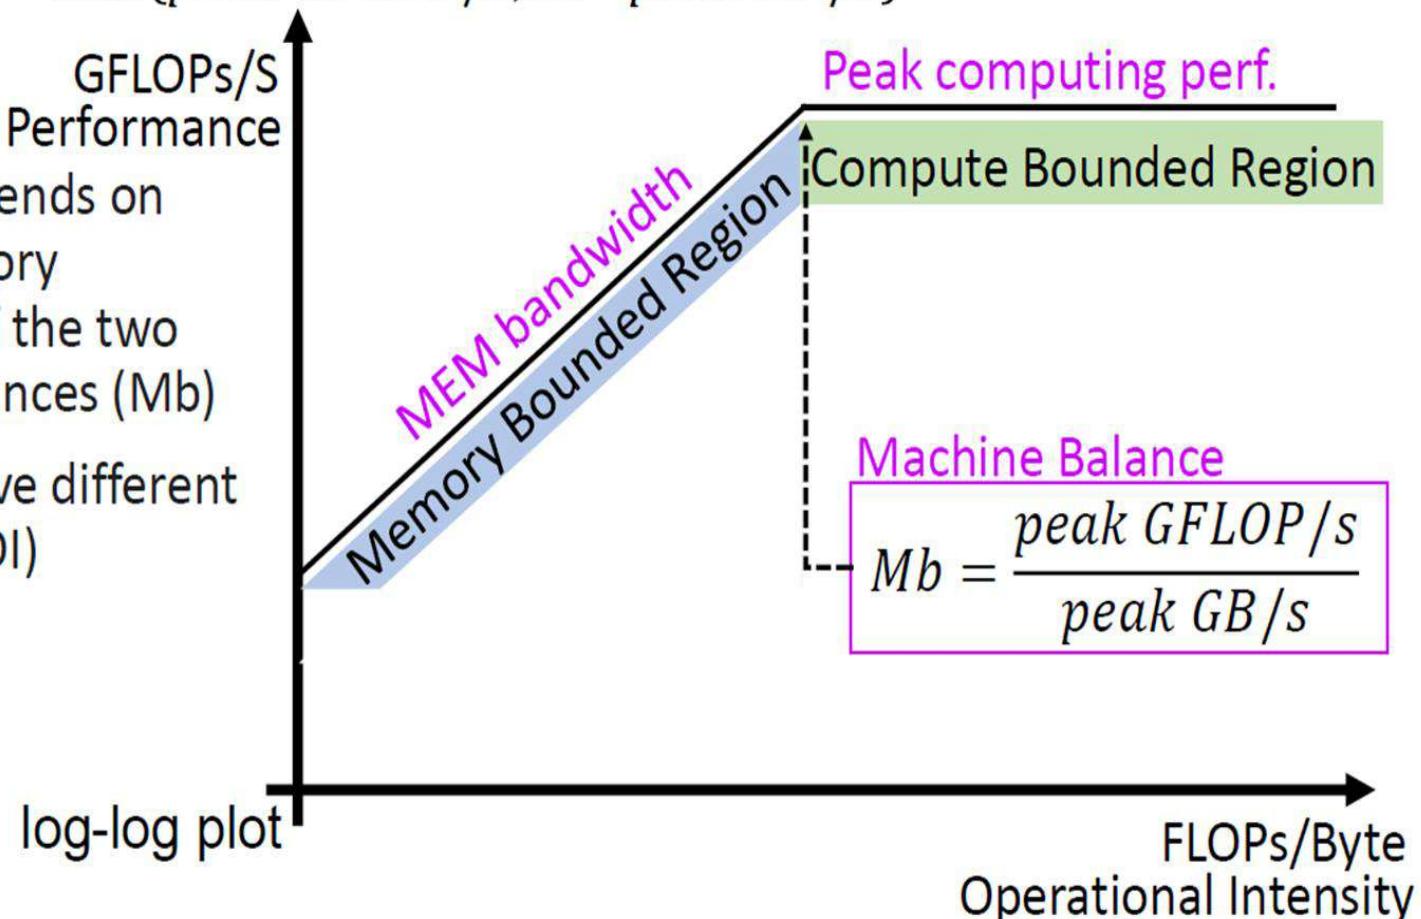
# Compute and Memory Performance

## Roofline Model:

$$\text{Attainable GFLOP/s} = \min\{\text{peak GFLOP/s}, OI \cdot \text{peak GB/s}\}$$

- System Performance depends on both Compute and Memory Performance, the ratio of the two defines the Machine Balances (Mb)
- Different Applications have different Operational Intensities (OI)

$$OI = \frac{\text{Num of FLOP}}{\text{Num of Byte}}$$



Use the scaling function G(pM) in memory-bound as the Operational Intensities (OI)

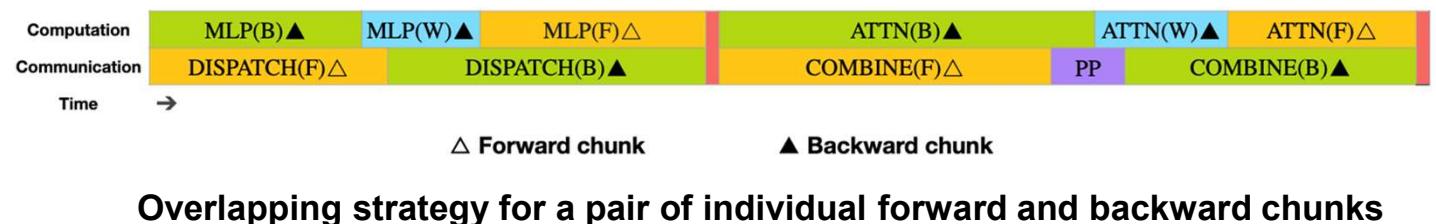
Micron E. Confalonieri, IEDM Short Course 2024

**Memory-Bound Model Impact**



# The deepseek Example: data-access optimization

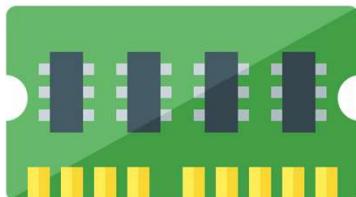
- DeepSeek MoE uses smaller, fine-grained experts to enhance data parallelism
- MoE Expert Routing ensures load balance and prevents parallelism inefficiencies
- Optimized pipeline & overlapping hide data movement latency
- All-to-All Communication Optimizations mitigate the data movement overheads
- Memory-Efficient Training & FP8 Mixed Precision reduce memory requirements and mitigate data movement pressure
- Still: DeepSeek inference spends **12x more time** on data access than of computing time



DeepSeek-V3 Technical Report



# (Performance) Bound is not just Capacity



## Just Buy More DRAM



DRAM is **very** expensive



DRAM has a  
**very** high  
energy cost

Only **20%** of memory  
used in 20% of time

The **Memory-wall**  
problem

Larger memory means  
**longer** searching time

	Price	Capacity	\$/GB
DRAM	\$106	32GB	\$3.30
PMEM	\$275	128GB	\$2.10
NVMe	\$169	2TB	\$0.08
SATA SSD	\$175	4TB	\$0.04
HDD	\$369	18TB	\$0.02

**40x more  
expensive  
than NVMe**

**Simply increasing DRAM is  
not sustainable!**



# Existing Solution: ASIC from CPU side

## ■ GPU, DSP, AI Chip

- GPU is a chip tailored to graphics processing, DSP is for signal processing, and AI chip is designed to do AI tasks



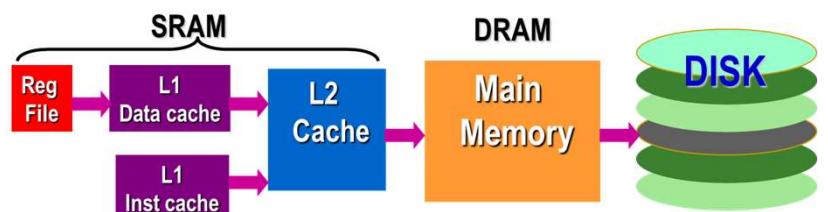
## ■ Limited solution

- Assume data are on the chip



## ■ Limited application

- *Computation Accelerator*
- Please recall our memory-bound results for multicore



The Traditional Approach: Memory Hierarchy



# New Solution: PIM chip

## ■ PIM

- ❑ Processing in memory (also called processor in memory) is the integration of a processor with RAM on a single chip.
- ❑ NDP (Near-memory Data Processing)
- ❑ ISP (In-Storage Processing)



## ■ Computer power is weak

- ❑ A full kitchen needs a refrigerator



## ■ Limited application

- ❑ *Data movement reducer*
- ❑ A helper/mitigator

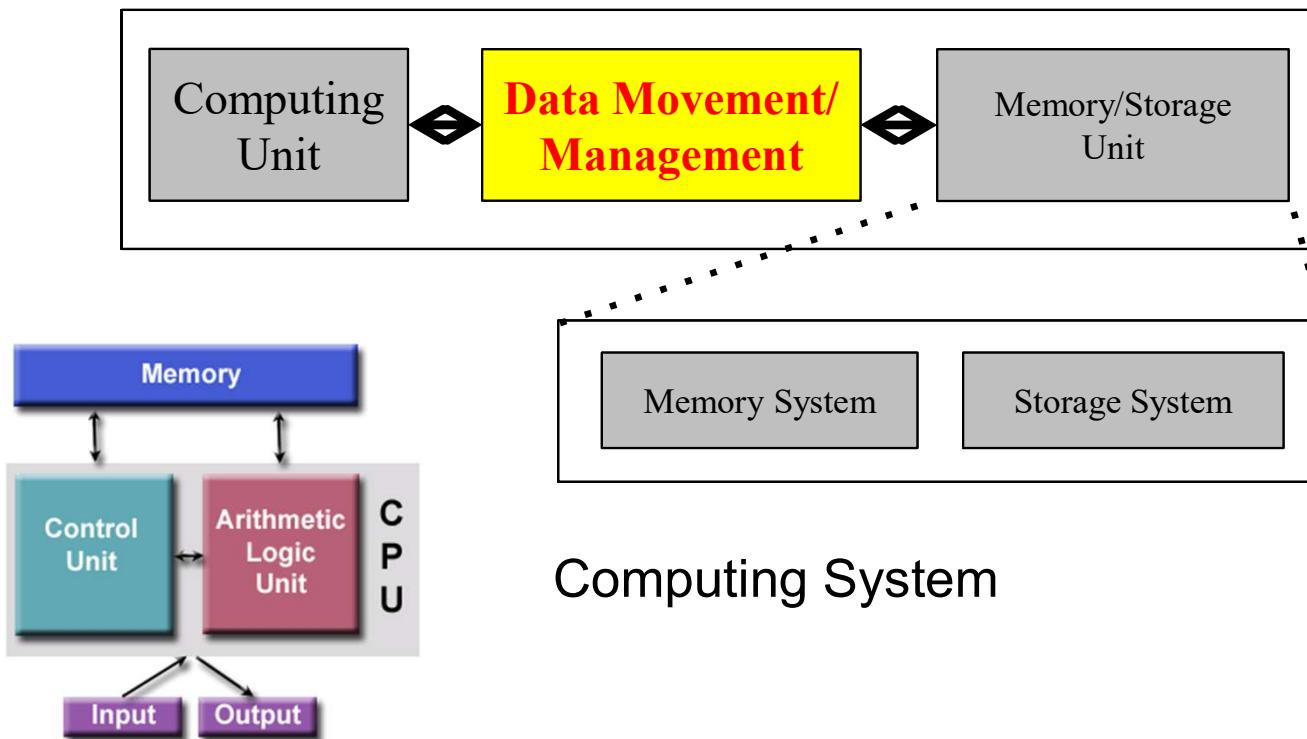
**How to use it?**





# Our Solution: A new component of von Neumann

- Can we make the von Neumann machine **data efficient**?
- **Yes**: focus on data and data access delay (data centric)
- **How**: *Advancing current memory-wall solutions*



**The Problem is in the Middle**

**Research**

- ❖ System Research
- ❖ Big Data/AI Systems
- ❖ High Performance Computing Data Systems

**Data Centric Imperative**

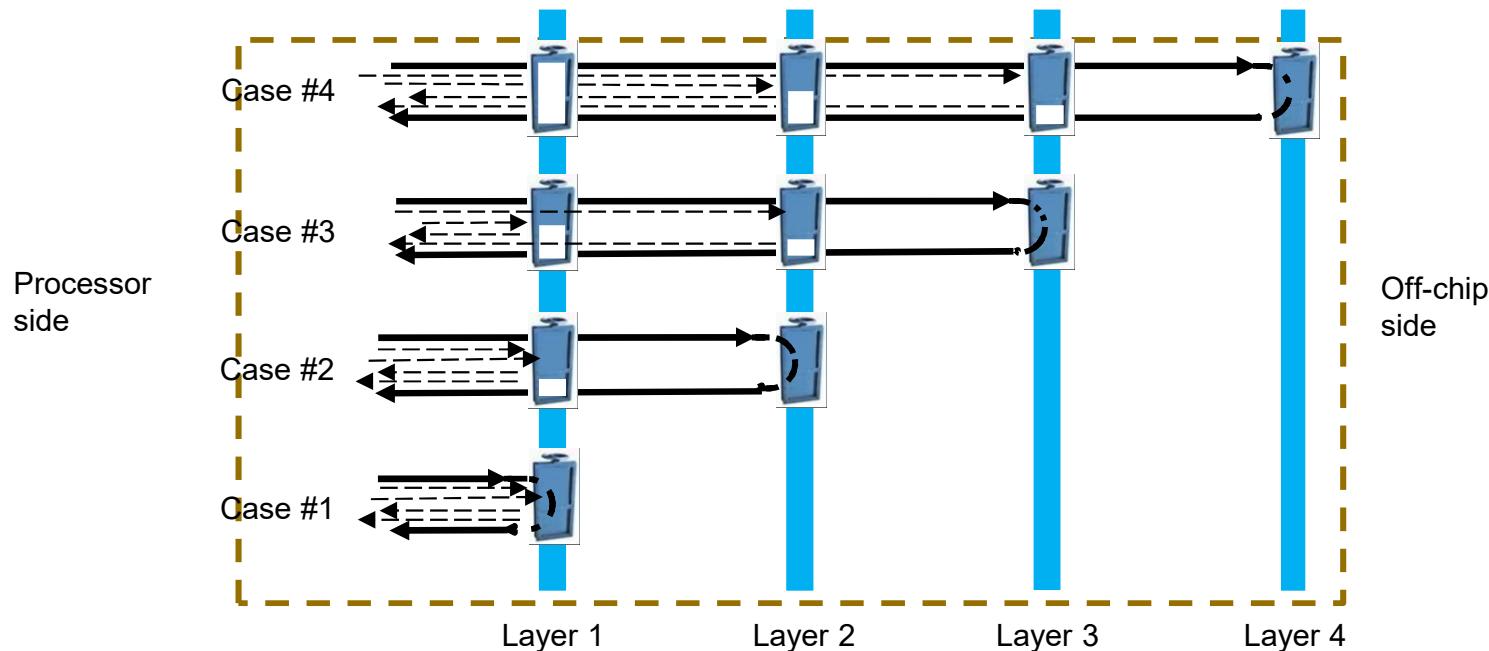


# Dataflow under the von Neumann Machine

## Memory Stall Time (MST)

$$\text{CPU.time} = \text{IC} \times (\text{CPI}_{\text{exe}} + \text{Memory stall time}) \times \text{Cycle.time}$$

*Reduce the Memory Stall Time to minimum*



(including in place computing)

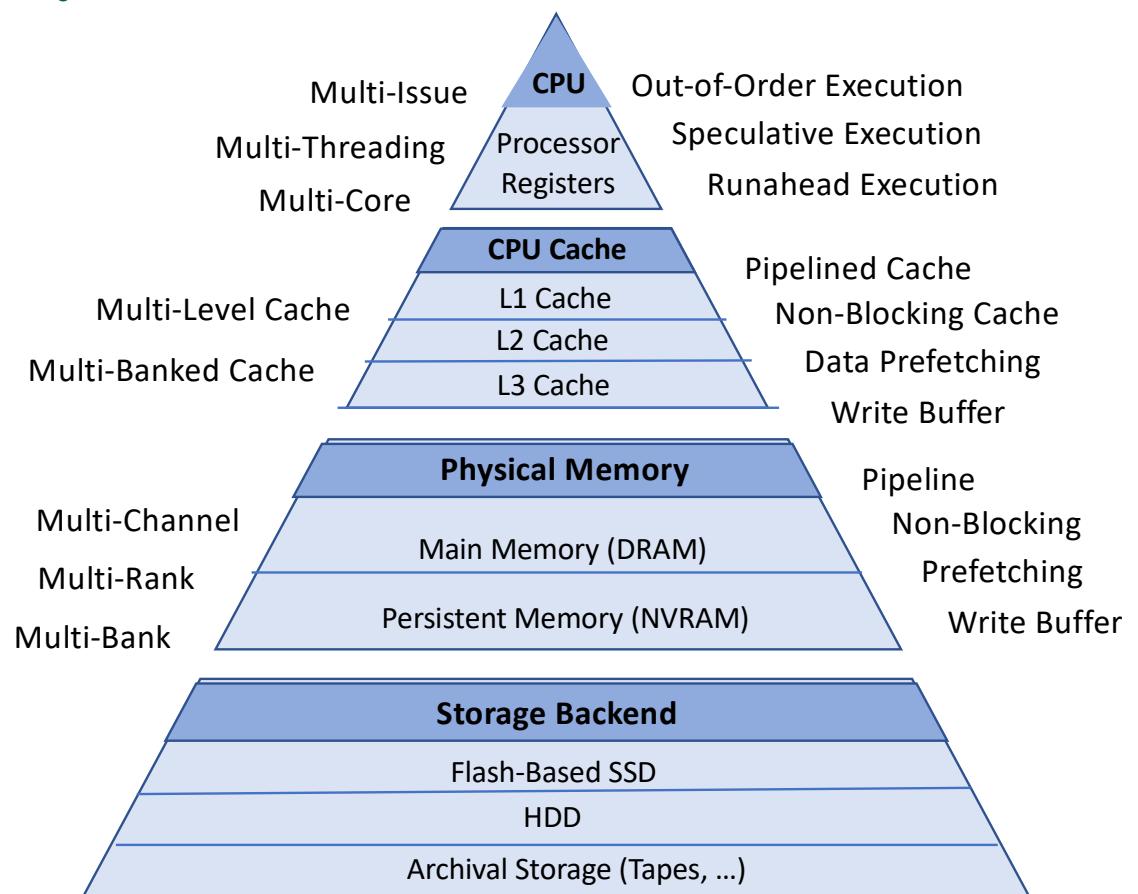
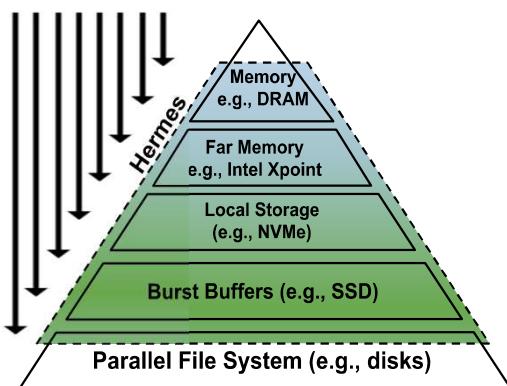
**Utilize current memory system**



# New Disruptive under Existing System: Integration & Concurrency

- Advanced solution: Deep Memory/IO Hierarchy
  - New technologies and complexities
  - **Locality** and **Concurrency**

*Our Idea:*  
**Integrated Memory System via  
Data Access Concurrency**



**Current Deep Memory-Storage Hierarchy  
with Concurrency**

**Challenges:** Currency, Integration, environment, technology, etc.

**Our Idea**



# C-AMAT: Concurrent Data Access fundamental

$$AMAT = H_1 + MR_1 \times AMP_1$$

Where  $AMP_1 = (H_2 + MR_2 \times AMP_2)$

- C-AMAT is Recursive

Where:

$$C-AMAT_1 = \frac{H_1}{C_{H_1}} + MR_1 \times \kappa_1 \times C-AMAT_2$$
$$C-AMAT_2 = \frac{H_2}{C_{H_2}} + MR_2 \times \kappa_2 \times C-AMAT_3$$

- H: the hit time
- MR: the miss ratio
- $C_H$  : the hit concurrency
- $\kappa$ : the overlapping ratio (**pure miss cycles** over miss cycles)
- A pure miss cycle is a miss cycle with no hit

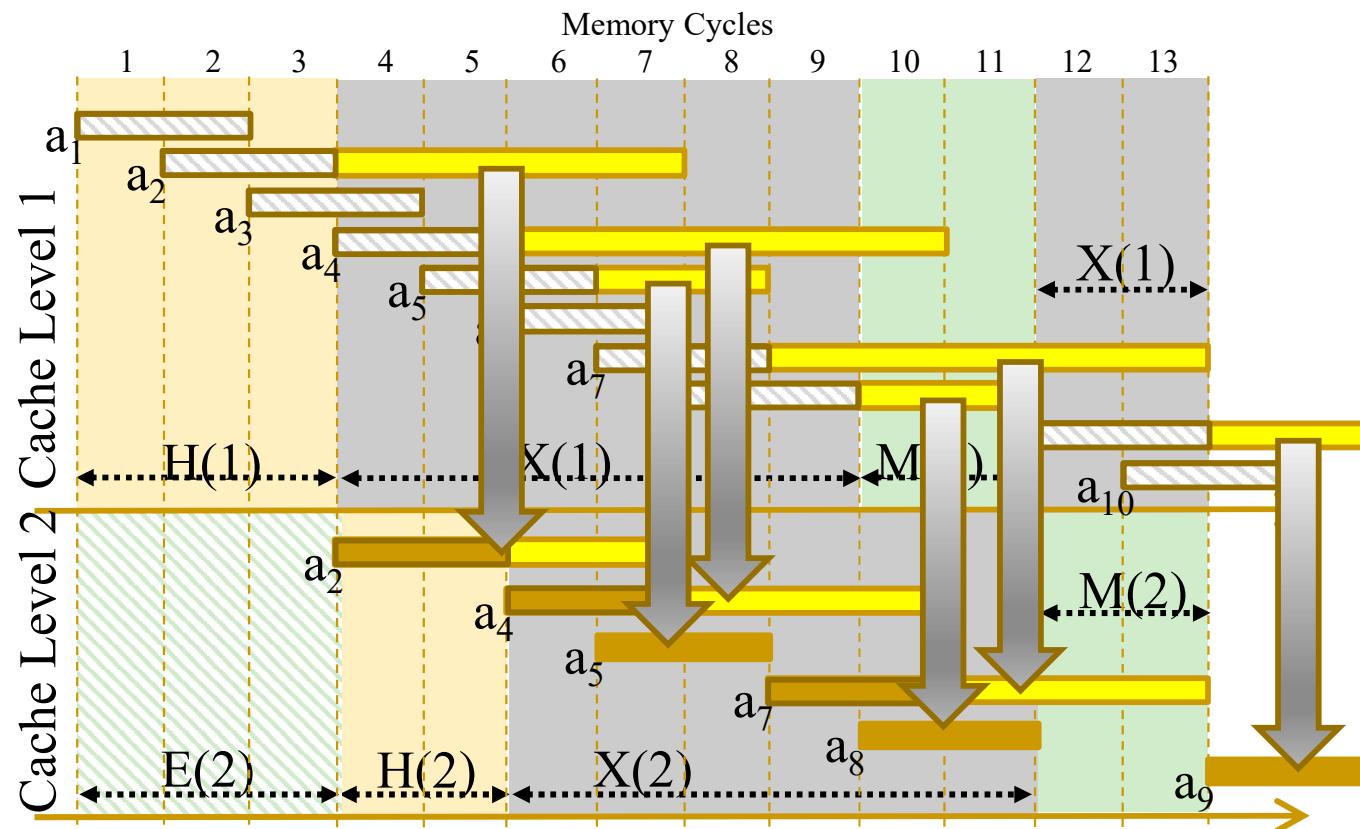
$$\kappa_1 = \frac{pMR_1}{MR_1} \times \frac{pAMP_1}{AMP_1} \times \frac{C_{m\_1}}{C_{M\_1}}$$



**Concurrence & locality are equally important**

Sun, Xian-He, and Dawei Wang. "Concurrent average memory access time." IEEE Computer 47, no. 5 (2014): 74-80.

## Theoretical Foundation (model)



L1 pure misses cause CPU stall

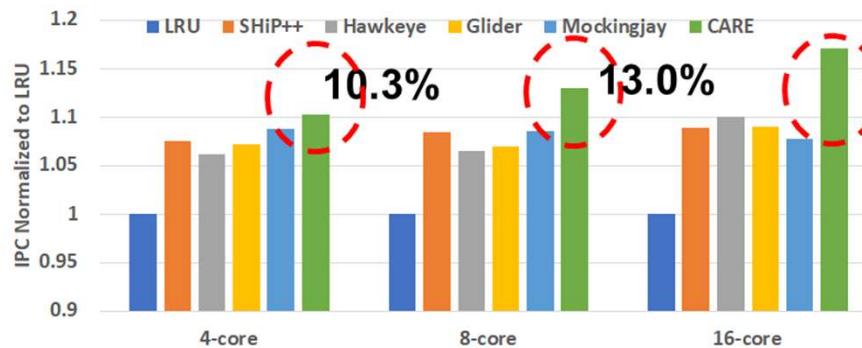
From concurrence-supported locality  
To Concurrence-aware Locality, & Locality-aware concurrence,



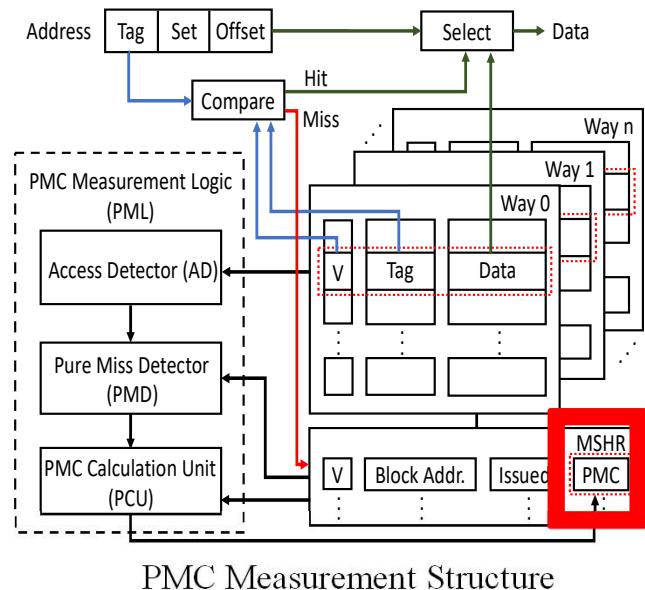
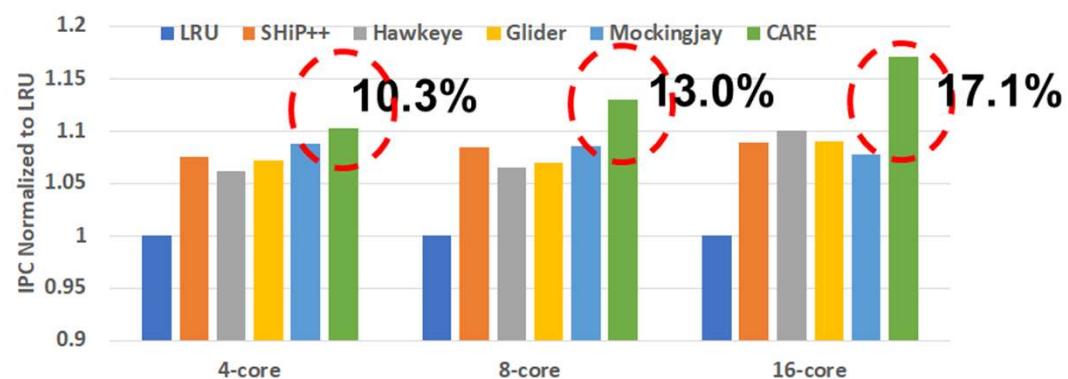
# Put in Action: Concurrency-Aware Cache Replacement

## Design

- Pure Miss Contribution (PMC)**
  - Pure miss is more important than miss
  - Pure miss has weight (pure miss cycles)
  - PMC is the number of pure miss cycles
  - PMC can be measured by monitoring MSHR occupancy
- PMC is Predictable**
  - PMC values of the misses caused by the same Program Counter (PC) are relatively stable



## SPEC & GAP Workloads

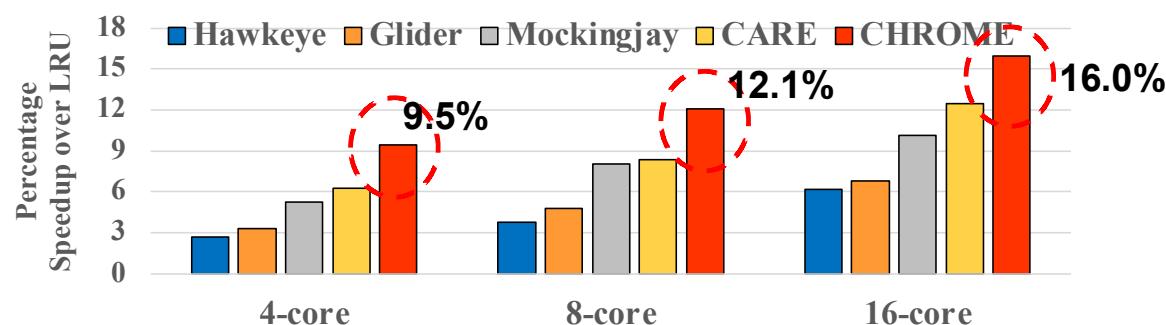
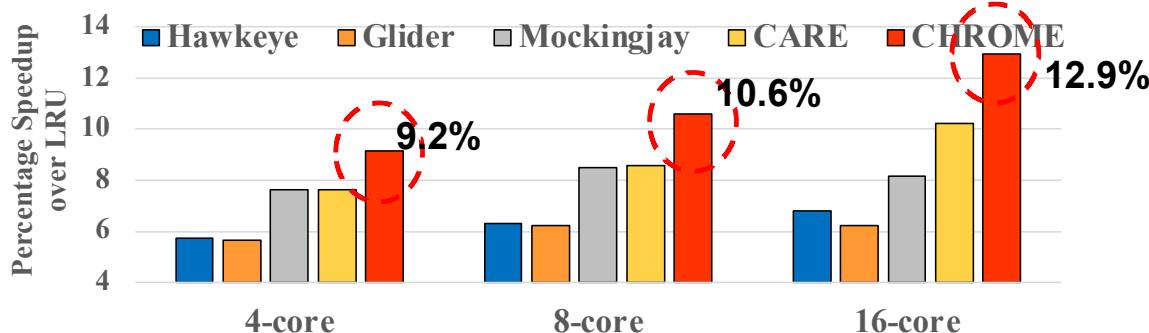




# Action: Concurrency-Aware Holistic Cache Management

## CHROME:

- Integrates **cache bypassing** and **replacement** with pattern-based **prefetching**
- Dynamic Online Learning**: for varying workloads and system configurations
- Multiple Program Features** for understanding of memory access patterns
- Concurrency-Aware Rewards for concurrent accesses (C-AMAT model)**



**CHROME outperforms state-of-the-art cache management schemes**

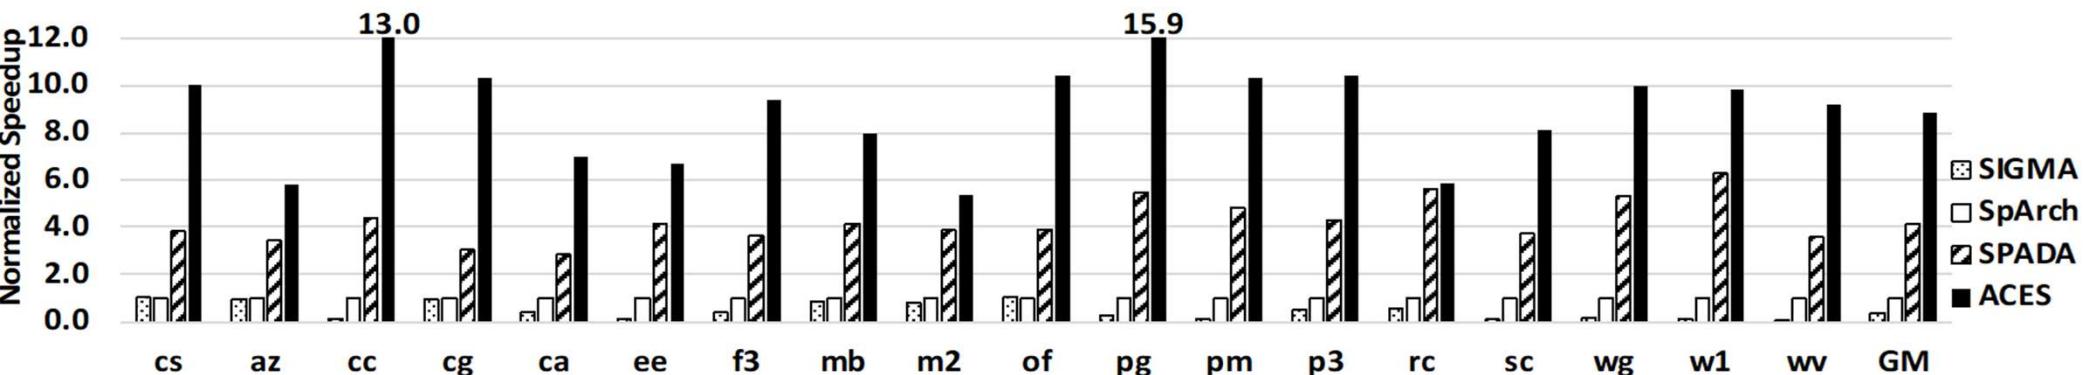
- Effective across both SPEC and GAP memory-intensive workloads
- Demonstrates strong scalability



## Action: Concurrency-Aware SpMM Accelerator

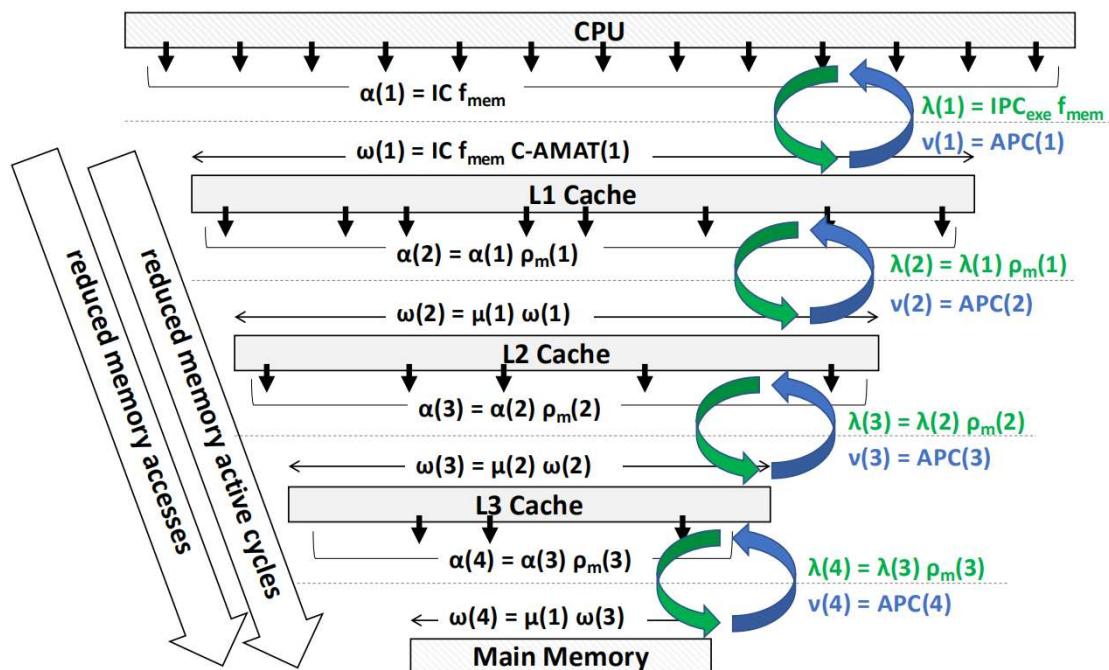
### ACES : Concurrency-Aware SpMM Accelerator

- **Adaptive Execution Flow:** better software concurrency
- **Non-Blocking (NB) Buffer:** better hardware concurrency
- **Concurrency-Aware Cache Replacement:** Considers **Reuse Distance (RD)** and **Fiber Density (FD)**, allows all cache lines of a row to be accessed **concurrently** without any misses



- ACES consistently provides **optimal** performance across **all** workloads
- **25.5×** over SIGMA, **8.9×** over SpArch, and **2.1×** over SPADA
- Non-Blocking (NB) Buffers enable greater concurrency support
- Concurrency-aware cache replacement policies further exploit this hardware-supported parallelism

- If each layer's performance is matched, then we get the top layer's performance and the lowest layer's capacity
  - The *Matching* ratio values of request and supply at each layer are given and the matching process is well designed & analyzed



## **Simulatable → Measurable → Controllable → Optimizable**

Y. Liu and X. Sun, "LPM: A Systematic Methodology for Concurrent Data Access Pattern Optimization from a Matching Perspective," in IEEE Transactions on Parallel and Distributed Systems, vol. 30, no. 11, pp. 2478-2493, 1 Nov. 2019.

# Theoretical Foundation (modeling)



# A Matching Example (increase performance)

$$LPMR = \frac{IPC_{exe} \times f_{mem}}{APC} \quad (1)$$

- **Assume:**

$$Cycle_{CPU} = 2 \text{ ns}$$

$$Cycle_{mem} = 8 \text{ ns}$$

$$f_{mem} = 20\%$$

$$IPC_{exe} = 2.5$$

$$APC = 1$$

- **Then:**

- In 8 ns, there is 1 memory cycle, and the data supply rate is:

$$APC * N_{Cycle_{mem}} = 1$$

- In 8 ns, there is 4 cpu cycle, and the data request rate is:

$$IPC_{exe} * N_{Cycle_{CPU}} * f_{mem} = 2$$

- **So:**

- The data supply rate does not match the data request rate. We can improve memory *concurrency* by adding memory banks or ports to increase data supply ability to adjust the data supply rate. For example, increase APC from 1 to 2, so:

$$APC * N_{Cycle_{mem}} = 2$$

- And the data supply rate matches with the data request rate.

## Case Study



# Matching from Another Angle (power consumption)

$$LPMR = \frac{IPC_{exe} \times f_{mem}}{APC} \quad (1)$$

- **Assume:**

$$Cycle_{CPU} = 2 \text{ ns}$$

$$Cycle_{mem} = 8 \text{ ns}$$

$$f_{mem} = 20\%$$

$$IPC_{exe} = 2.5$$

$$APC = 1$$

- **Then:**

- In 8 ns, there is 1 memory cycle, and the data supply rate is:

$$APC * N_{Cycle_{mem}} = 1$$

- In 8 ns, there is 4 cpu cycle, and the data request rate is:

$$IPC_{exe} * N_{Cycle_{CPU}} * f_{mem} = 2$$

- **So:**

- The data supply rate does not match the data request rate. We can decrease the CPU *frequency* (increase CPU cycle) to adjust the data request rate. For example, increase the CPU cycle from 2 ns to 4 ns, the  $N_{Cycle_{CPU}}$  would be 2, so:

$$IPC_{exe} * N_{Cycle_{CPU}} * f_{mem} = 1$$

- And the data request rate matches with the data supply rate.

## Case Study



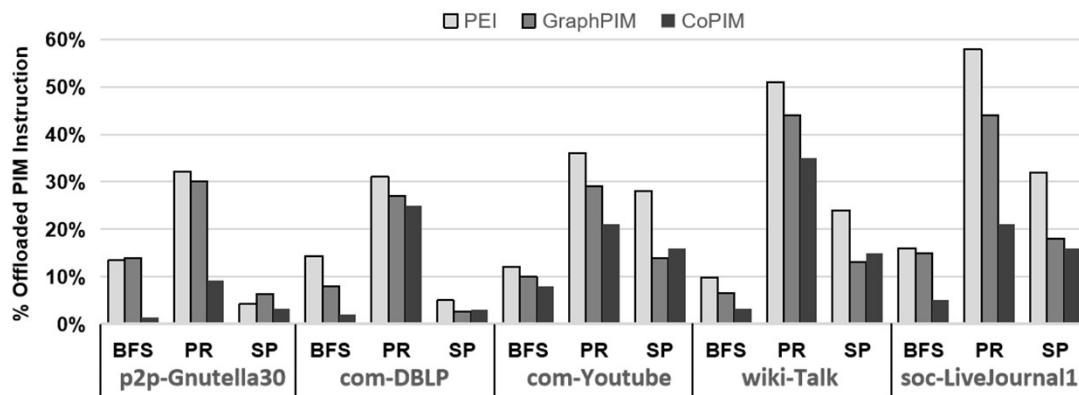
# Case Study: Put PIM into the Formula

## Memory Stall Time (MST)

$$\text{CPU.time} = \text{IC} \times (\text{CPI}_{\text{exe}} + \text{Memory stall time}) \times \text{Cycle.time}$$

$$\begin{aligned}\text{CPU.time} = & \text{IC}_{\text{exe}} \times (\text{CPI}_{\text{exe}} + \text{Memory stall time}) \times \text{Cycle.time} \\ & + \text{IC}_{\text{pim}} \times (\text{CPI}_{\text{pim}} + \text{PIM stall time}) \times \text{Cycle.time}_{\text{pim}} - \text{Overlapping}\end{aligned}$$

- PIM is a way to reduce request and is a trade-off of computing and MST
- Use Pure Miss as a indicator for offloading
- Result: data movement cost **decides** where to do the computing



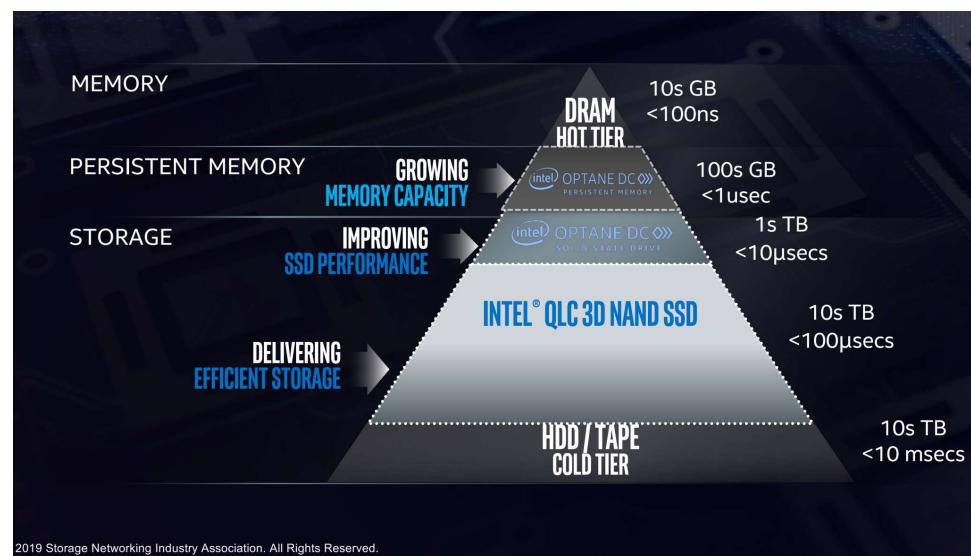
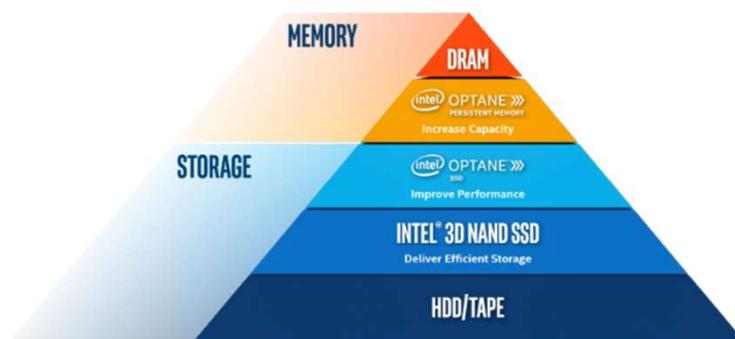
Percentage of offloaded instructions into memory

Speedup by  
**19.5%** than **PEI**  
with **51.1%** fewer  
offloaded  
instructions



# Case Study 4: The Intel Optane Technology

- Manufacture hardware integration (matching)
- Challenges
  - Data access is application dependent
  - Hardware **elasticity**
  - Global optimization
  - OS, etc.
- *An architecture ahead of time?*

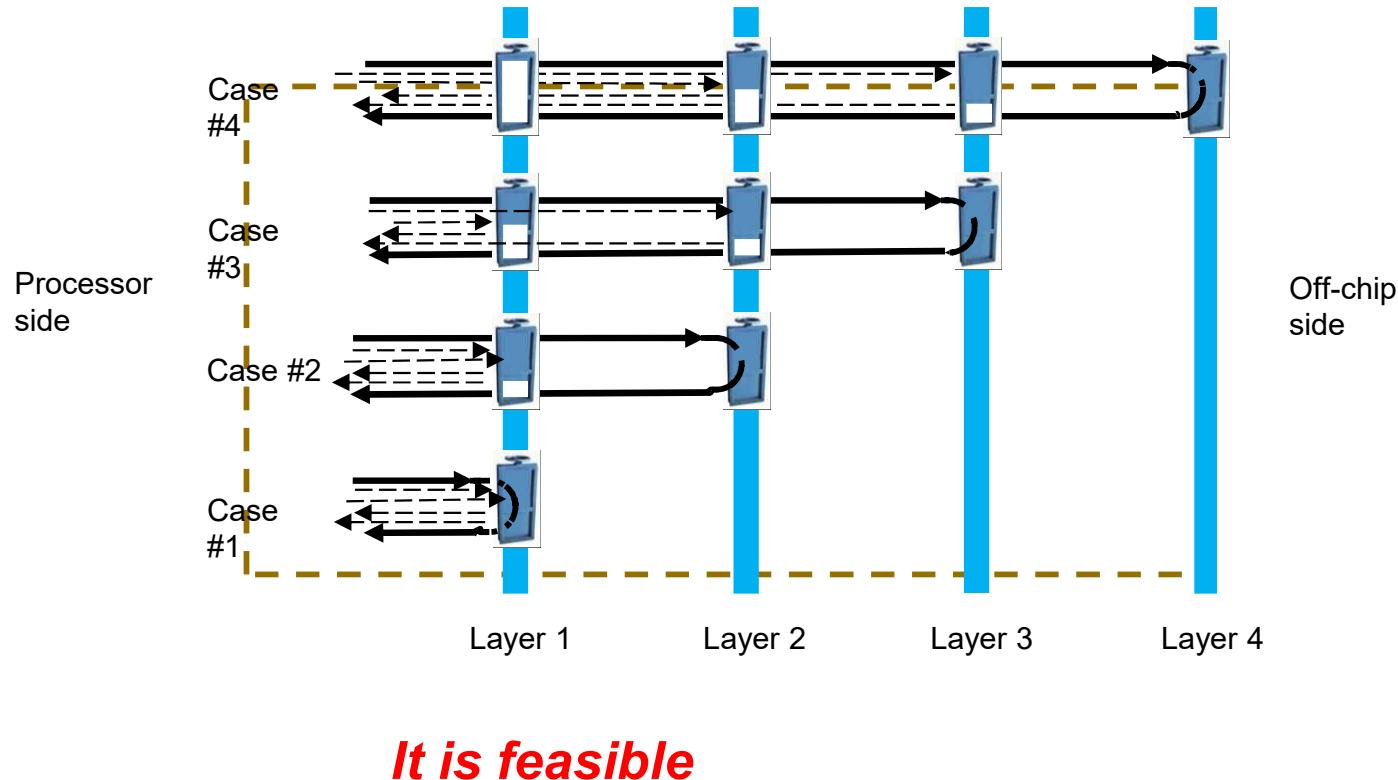


N. Zhang, B. Toonen, X-H. Sun, B. Allcock, "Performance Modeling and Evaluation of a Production Disaggregated Memory System," International Symposium on Memory Systems (MEMSYS'20), Sept. 2020

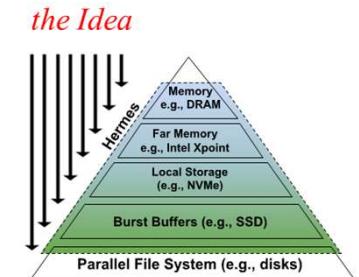


# Dataflow under the von Neumann Machine

- Use **LPM** to utilize the memory system performance
  - Identify the number of layers and carry matching at each layer
- Use **C-AMAT** (concurrency) and others to perform the matching
- Goal: ***Dataflow*** from the source to compute unite with minimum stall



***It is feasible***





# The *Dataflow*<sub>V</sub> Complexity

## Possible Ways to Match

- Reduce request
  - Improve locality, **processing-in-memory (PIM) & in place computing**, algorithm improvement, data movement reducer, etc.
- Improve supply
  - Improve data access concurrency , buffer, new technology, etc.
- Mask the difference
  - **Overlapping** computing with data access delay (**pure miss**), prefetch, etc.
- A more balanced design, etc.

Hardware technology, compiler technology, application algorithm designs, system scheduling, system support, co-design, integration, etc.

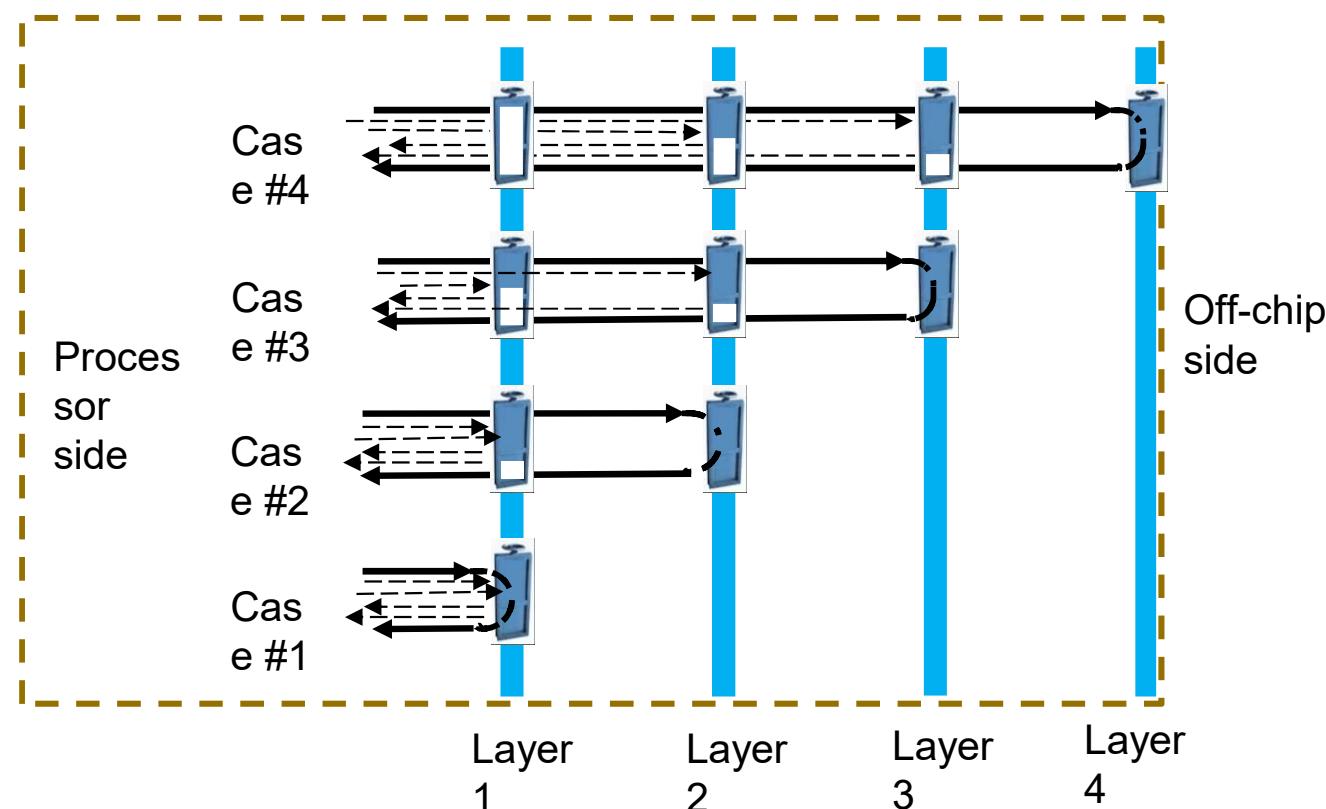
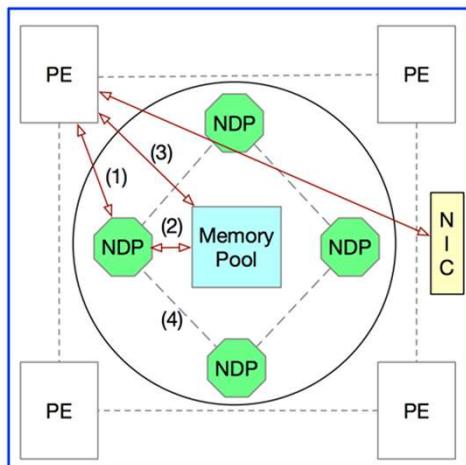


# Action: Towards a Unified Memory-centric Computing System with Cross-layer Support



No delay data access: An integration of memory, storage, architecture (PIM, NDP, CXL, etc), OS, compiler, runtime systems

- Is your new device/technology helpful?
- How should they be used together?



New advances in computer architecture

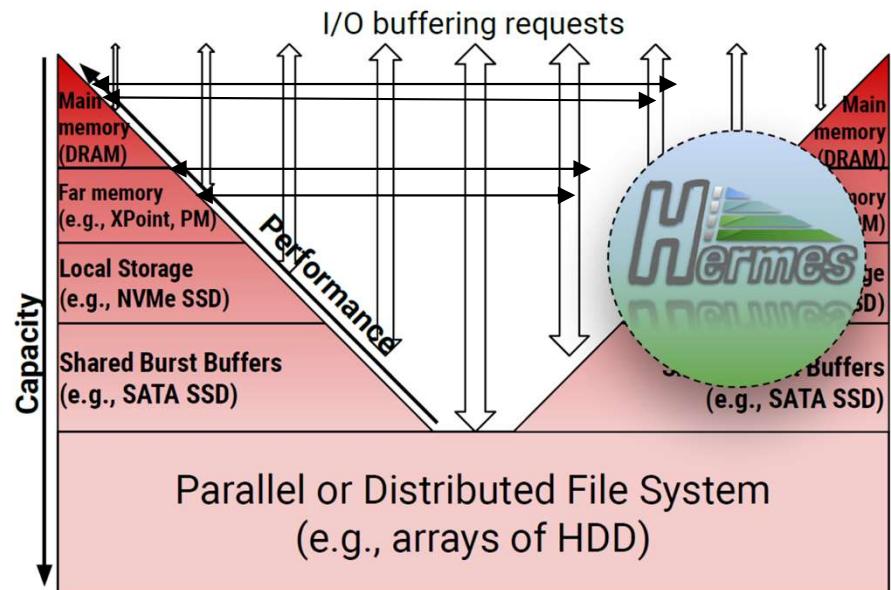
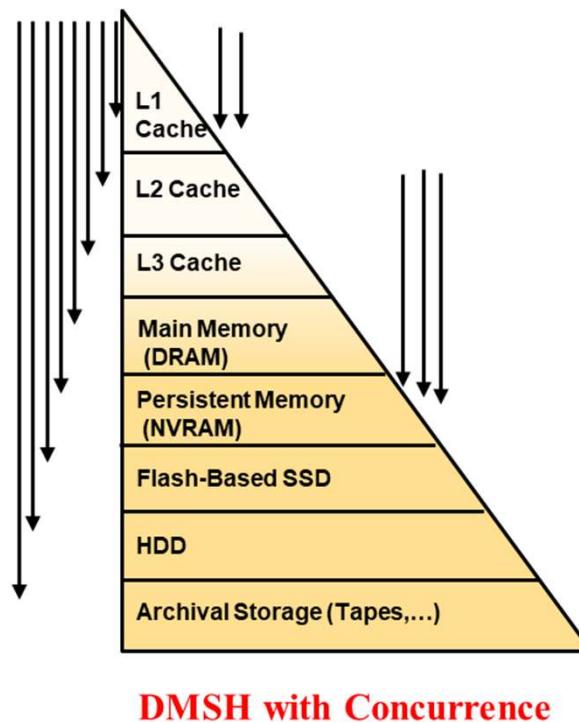
Dataflow under the von Neumann machine



# Put in Action: Hermes: A Smart Multi-tiered I/O Buffering System



- Storage as a layer of the memory hierarchy
- Selective cache, data concurrency, matching
- Independent management of **each tier**
- Optimized Access Latency, Throughput, Capacity, Network topology, etc.



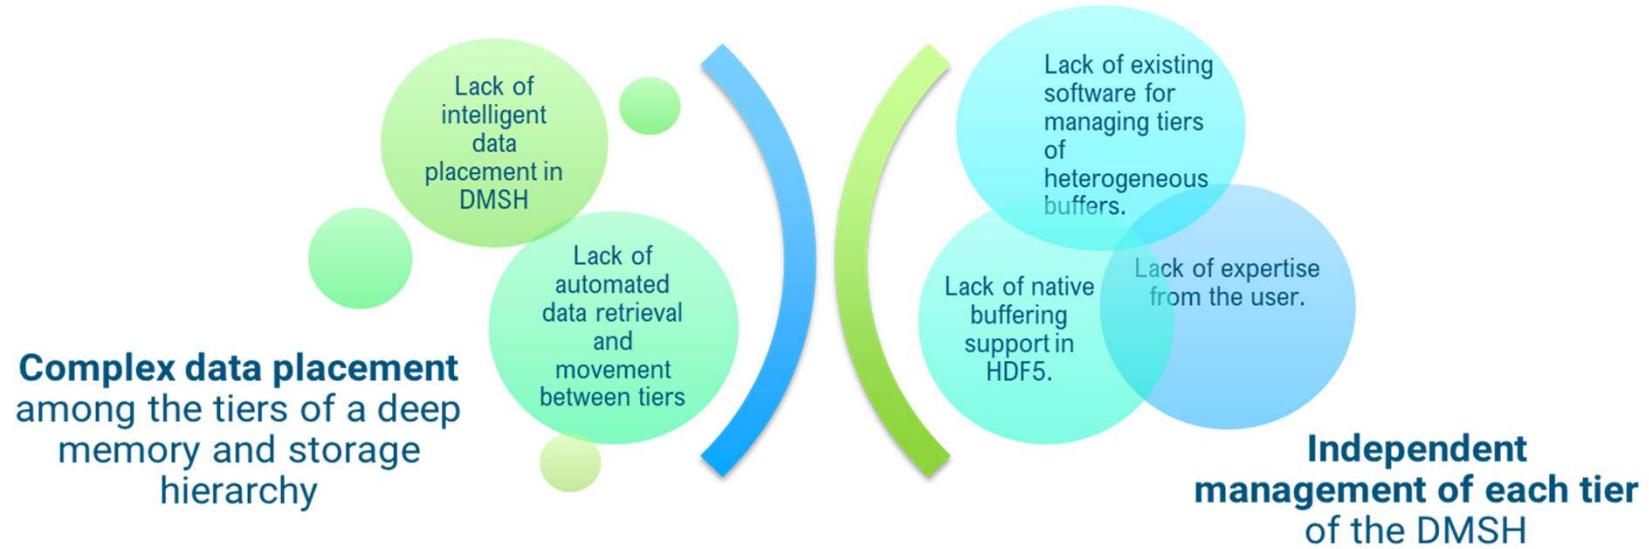
A. Kougkas, H. Devarajan, and X.-H. Sun, "I/O Acceleration via Multi-Tiered Data Buffering and Prefetching," Journal of Computer Science and Technology, vol. 35, no. 1, pp. 92-120, Jan. 2020

## Our First Data Management System



# The Development Challenges

- Application-aware multi-tier optimization
- Complexity in software development
- Complexity in decision making
- An example of memory/storage integration
- An implementation of the layered *matching and data flow* concept



A. Kougkas, H. Devarajan, and X.-H. Sun, “Bridging Storage Semantics using Data Labels and Asynchronous I/O,” ACM Transactions on Storage, Vol 16, No 4, 2020



# Technical Preparation

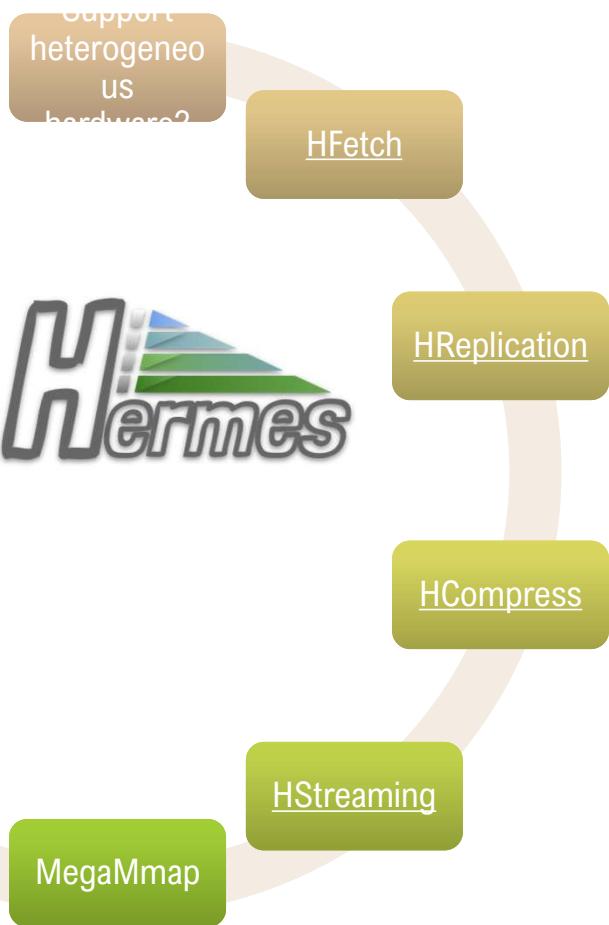
Hermes introduce several independent technologies

Explore our code at  
<https://github.com/grc-iit>.

[github.com/grc-iit/hfetch](https://github.com/grc-iit/hfetch)

[github.com/grc-iit/hreplica](https://github.com/grc-iit/hreplica)

[github.com/grc-iit/hcompress](https://github.com/grc-iit/hcompress)



**ML-driven Auto-Tuning and Auto-Configuration:** N. Rajesh, K. Bateman, S. Byna, J. L. Bez, A. Koukas, X.-H. Sun. "TunIO: An AI-powered Framework for Optimizing HPC I/O", IEEE IPDPS 2024.



# Other Hermes Tools and Services

- **Application orchestrator<sup>1</sup>:**
  - schedules application's access to Hermes in multi-tenant scenarios, mitigating interference by shared buffers.
- **System telemetry<sup>2</sup>:**
  - captures in real-time the status of tiers (e.g., load, remaining capacity, etc.,) to enable efficient data placement by Hermes.
- **Application I/O profiling<sup>3</sup>:**
  - classifies the I/O phases of applications to enable optimal Hermes configuration tuning based on incoming access patterns.
- **Hermes Container Library<sup>4</sup>:**
  - implements STL-like data structures (maps, list, etc) that can be distributed across nodes and within available storage tiers.
- **Resource Management<sup>5</sup>:**
  - provision and allocate the best storage hardware resources available based on application demands and QoS guarantees
- **The Memory Management System**

1.A. Kougkas, H. Devarajan, X-H. Sun, and J. Lofstead. "Harmonia: An Interference-Aware Dynamic I/O Scheduler", In IEEE International Conference on Cluster Computing 2018 (Cluster'18)

2.N. Rajesh, H. Devarajan, J. Cernuda Garcia, K. Bateman, L. Logan, J. Ye, A. Kougkas, and X-H. Sun. 2020. "Apollo: An ML-assisted Real-Time Storage Resource Observer". In 30th International Symposium on High-Performance Parallel and Distributed Computing (HPDC'21)

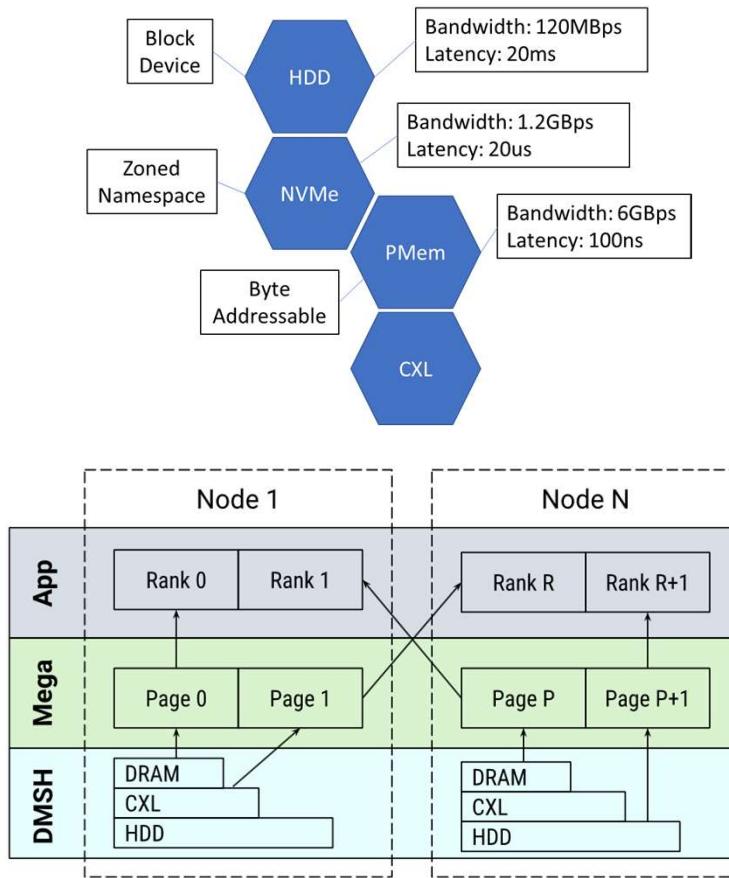
3.Hariharan Devarajan, Anthony Kougkas, P. Challa, Xian-He Sun, "Vidya: Performing Code-Block I/O Characterization for Data Access Optimization" In IEEE International Conference on High Performance Computing, Data, and Analytics 2018 (HiPC'18)

4.H. Devarajan, A. Kougkas, K. Bateman, and X-H Sun. "HCL: Distributing parallel data structures in extreme scales." In 2020 IEEE International Conference on Cluster Computing (Cluster'20).

5.K. Bateman, N. Rajesh, J. Cernuda, L. Logan, J. Ye, S. Herbein, A. Kougkas, and X.-H. Sun. "LuxIO: Intelligent Resource Provisioning and Auto-Configuration for Storage Services" In IEEE 29th International Conference on High Performance Computing, Data, and Analytics (HiPC'22)



# Blurring the Line of Memory & Storage



- Storage has emerged with similar performance to DRAM
- E.g., Compute Express Link (CXL) is emerging and allows storage to be accessed as memory
- **We can combine DRAM and storage into a tiered and portable distributed shared memory (DSM) abstraction**



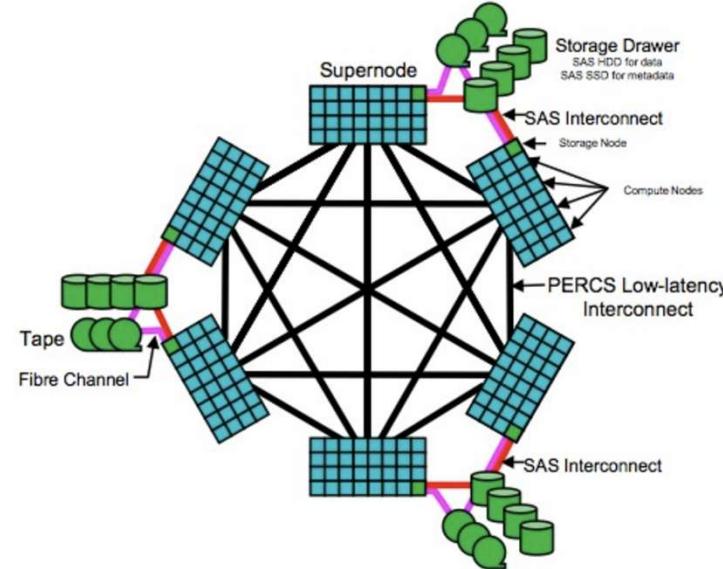
L. Logan, X.-H. Sun, A. Kougkas, "MegaMmap: Blurring the Boundary Between Memory and Storage for Data-Intensive Workloads", in the Proceedings of IEEE/ACM SC'24, Atlanta, USA, Nov. 17-22, 2024.



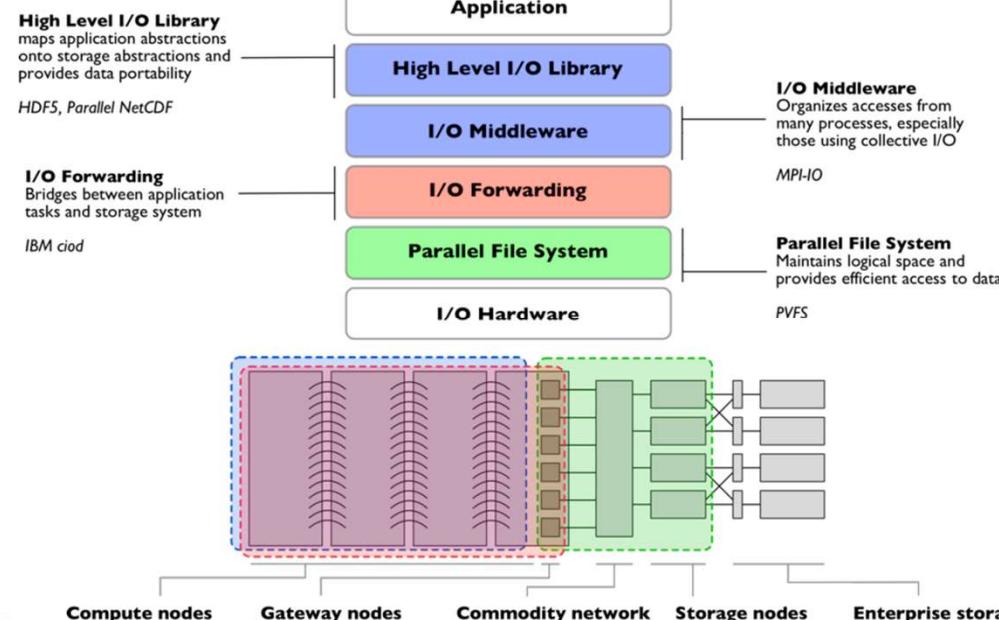
# The Complexity of HPC I/O Systems



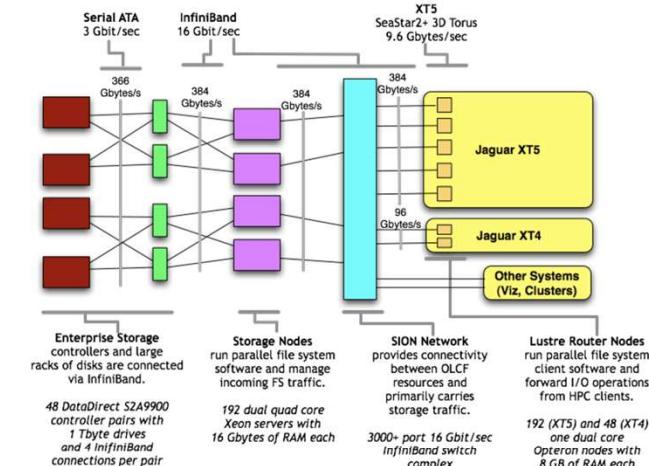
IBM Blue Gene/P system at Argonne National Laboratory.



Bluewaters Storage System



The I/O Software Stack of Intrepid

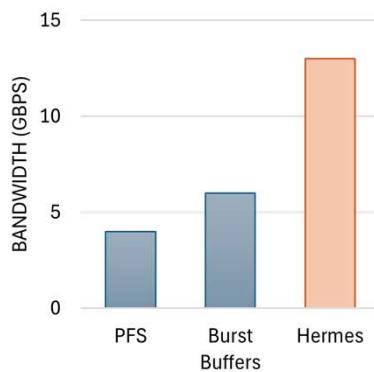


Jaguar Storage System

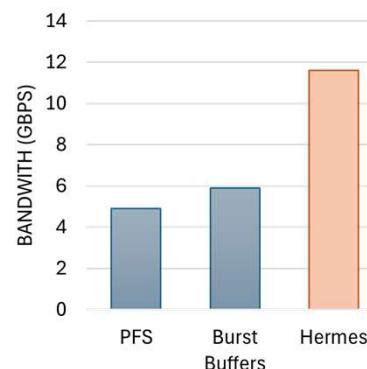
# Hermes System Results

Utilizing data access concurrence and memory hierarchy system with Multi-tiered I/O buffering

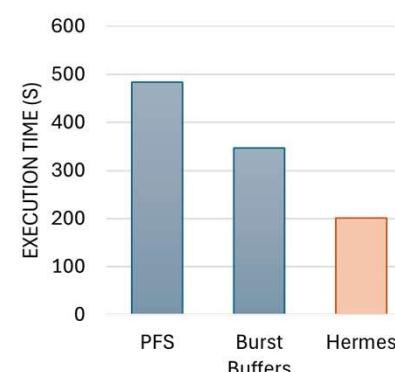
## HPC Simulations and Big Data Analytics



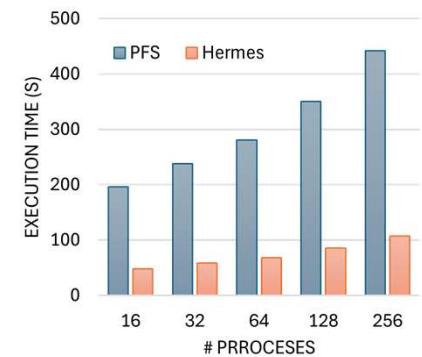
- Vector Particle-In-Cell (VPIC)
  - Write intensive simulation
  - **3x faster** than baseline (PFS)
  - 2.2x improvement compared to single-tier buffering solutions.



- Particle clustering (BD-CATS)
  - Read intensive analysis
  - 2.4x faster than baseline (PFS)
  - 2x improvement compared to single-tier buffering solutions.



- Gray-Scott reaction and diffusion models
  - Write intensive phase
  - **2.5x faster** than baseline (PFS)
  - 1.7x improvement compared to single-tier buffering solutions.

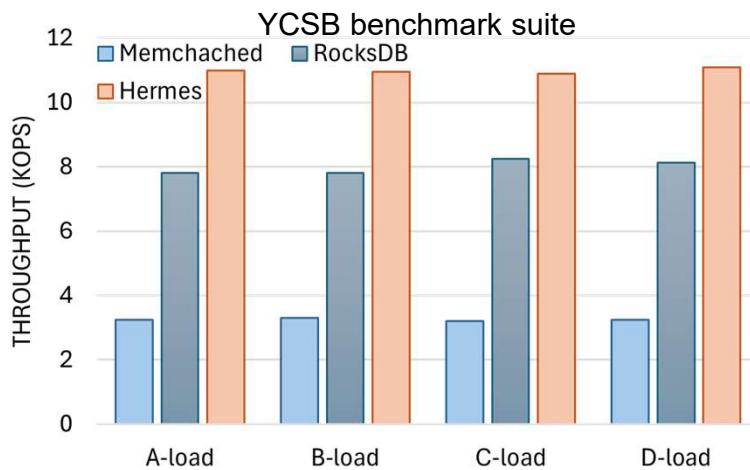


- Gray-Scott Analysis
  - Read intensive phase
  - 4x faster than baseline (PFS)

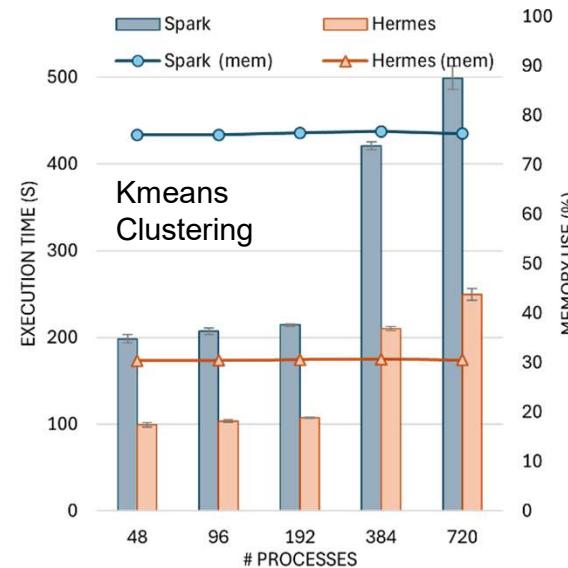
# Hermes System Results

Utilizing data access concurrence and memory hierarchy system with Multi-tiered I/O buffering

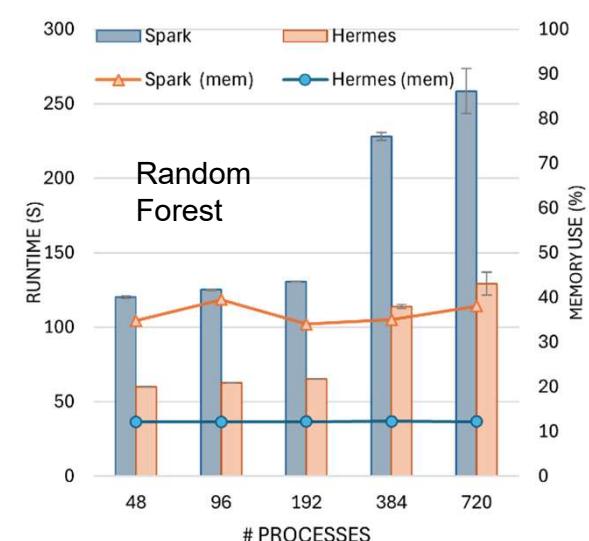
## Cloud and ML Workloads



- Performance superiority, using the same resources, for a variety of cloud workloads:
  - A. Webserver
  - B. Email Server
  - C. DB queries
  - D. KVS get/put



- Hermes outperforms SparkML in distributed ML algorithms by more than **3x highlighting** the performance benefits of data tiering.
- Hermes reduces DRAM usage by over **2x, enhancing data** processing capabilities without extra infrastructure overhead.





# After Hermes

- Application and Collaboration
  - Extended to Active (log) data: the **ChronoLog** project
  - Extended to Meta data: the Coeus project
  - Extended to workflow environment, the Cloud environment
  - Extended to AI applications
- Fundamental: Data centric system integration
  - *Unified Memory-centric Computing System (UMC2)*
    - *PIM, DPU, OS and compiler, memory/storage*
  - **LABIO**: a data-centric I/O system
    - LabStor and DTIO
- The next environment for AI: IO Warp
  - Bending the I/O Fabric for Advancing **AI-Infused Scientific Workflows**
    - StoreHub

New application  
New technology  
New environments

CCGrid 2021 The 21th IEEE/ACM International Symposium on Cluster, Cloud and Internet Computing (CCGrid 2021)

IEEE ACM INRIA University of Illinois at Urbana-Champaign

**Best Paper Award**  
**First Prize**  
DLIO: A Data-Centric Benchmark for Scientific Deep Learning Applications  
Presented to  
Hariharan Devarajan, Huihuo Zheng, Anthony Kougkas, Xian-He Sun and Venkatram Vishwanath

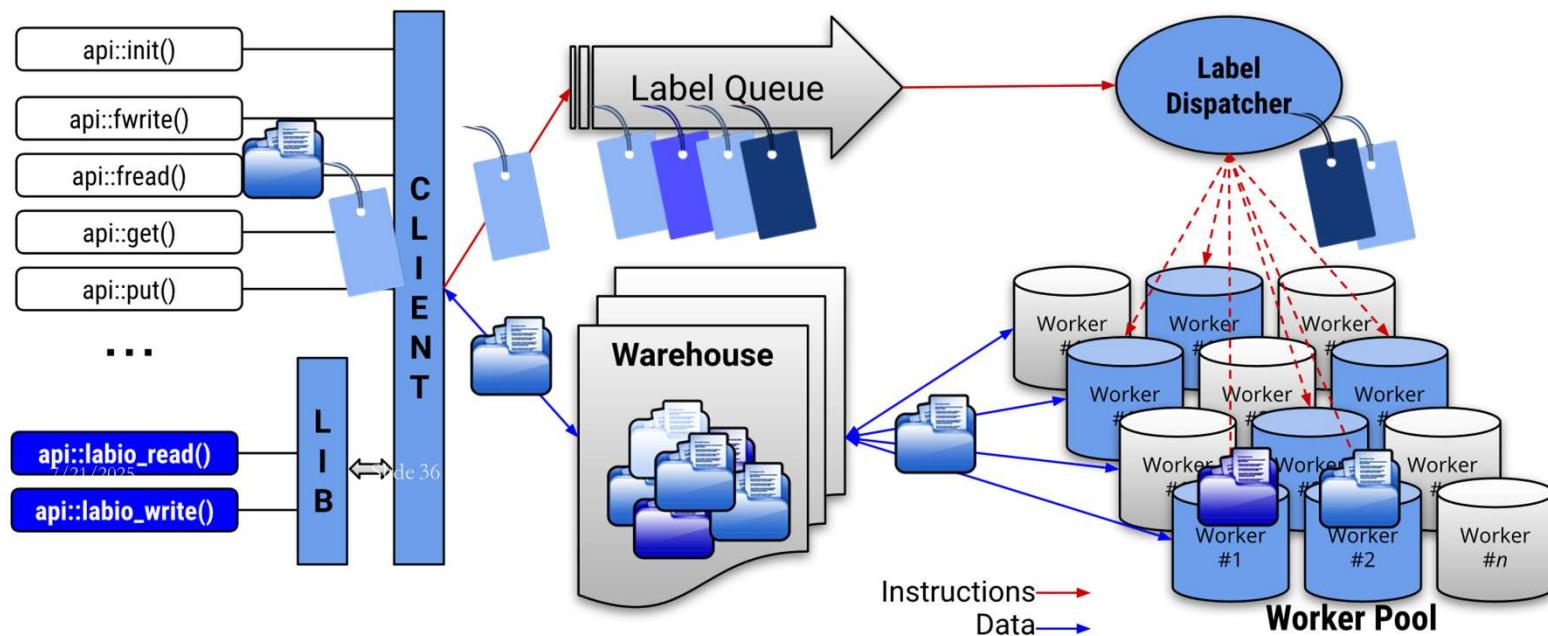
Laurent Lefevre, Program Chair, CCGrid 2021  
Stacy Patterson, Program Chair, CCGrid 2021  
Young Choon Lee, Program Chair, CCGrid 2021



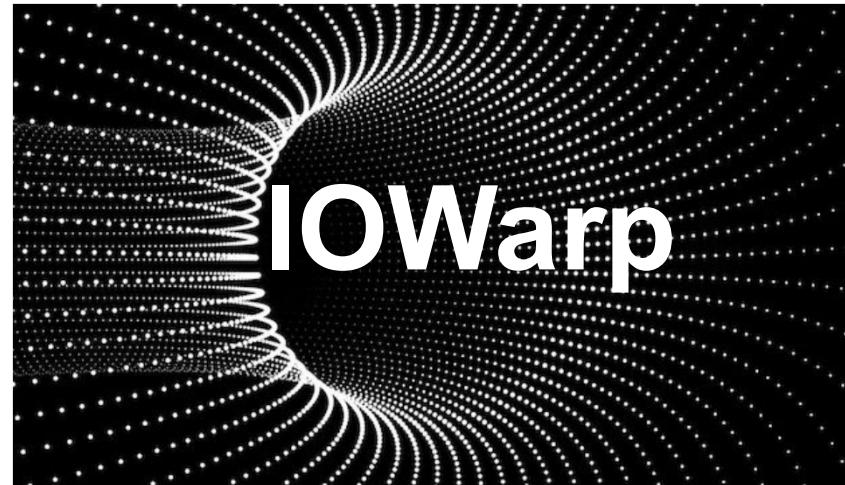
# LABIOS: Data Operation with Label



- Data requests are transformed into (data) **Label** units
  - A label is a tuple of an operation and a pointer to the data
- A dispatcher distributes labels to the workers
- Workers execute labels independently (i.e., fully decoupled)



A. Kougkas, H. Devarajan, J. Lofstead, X.-H. Sun; “*LABIOS: A Distributed Label-Based I/O System*”, in Proceedings of ACM HPDC ’19 (Best Paper Award)



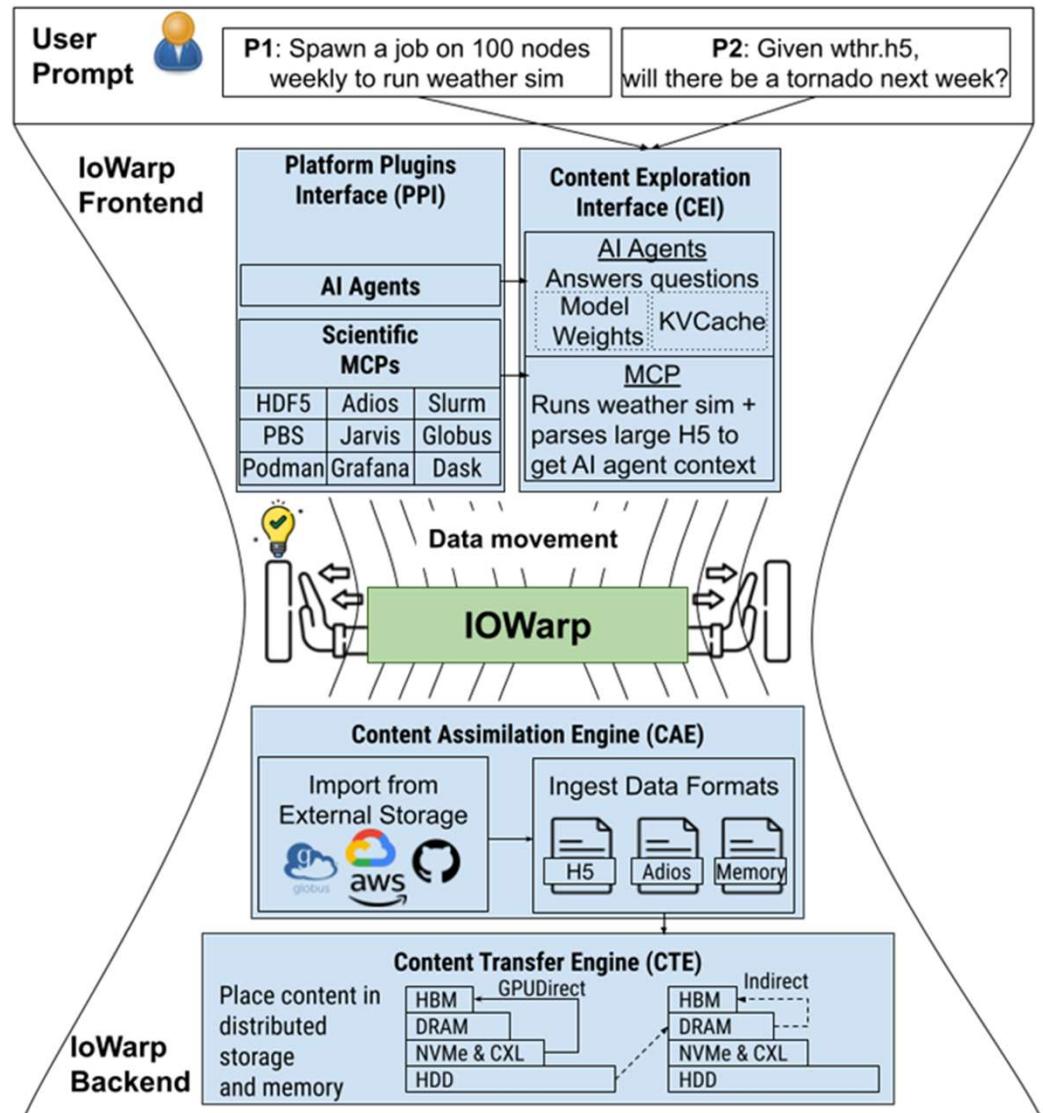
## IOWarp: The Continuation

IOWarp is a **data management** platform designed to address the unique challenges in big data and AI applications



# IoWarp: the Four Components

- Content Exploration
  - Comprehensive querying and indexing system and WarpGPT
- Platform Plugins
  - Connecting external tools and services
- Content Assimilation Engine
  - Convert diverse data formats into Content
- Content Transfer Engine
  - Manage efficient data transfer





# How Can You Get Involved?

The [IOWarp](#) organization is the primary contact point for all IOWarp content <https://github.com/iowarp>

Inside, new users can:

- Get started quickly with our [Installation Guides](#)  
<https://github.com/iowarp/iowarp-install>
- Explore the design and more complex features in the [Tutorial](#)

Our code bases welcome collaborators:

- Our [Runtime](#), modular execution engine for custom storage and data analytic pipelines.
- [Jarvis](#), a software deployment environment that abstract hardware and shell complexities.
- A collection of [Scientific MCPs](#) and an [Exploration Engine](#) enabling a natural language interface to scientific data.
- The [Assimilation Engine](#), a framework for transforming diverse format-specific data into an abstracted view called Content.
- The [Transfer Engine](#), a multi-tiered, dynamic, and distributed I/O buffering system that manages all data movement and placement.





## How Can You Get Involved?

# StoreHub

A Community Testbed





# Path Forward: A Community-Driven Initiative

- **Community-Centric Approach:** Identified needs through surveys, workshops (HDF User Group, PDSW, SSDBM), and direct conversations.  
<https://forms.gle/c6fmQkiSrCYWQMij9>
- **Build vital infrastructure** using community insights.
- **StoreHub** is the nation's premier *testbed for storage technologies*.  
<https://grc.iit.edu/research/projects/storehub/>
- **Pioneering Research:** Enabling breakthroughs in data management and storage research.
- **Broader Impact:** Enhancing education, workforce development, and fostering collaboration.





# Take Home Messages

---

- **Modeling is the foundation :**
  - From memory-bound to C-AMAT and LPM
  - Understanding and Solution
- **Solution:** Dataflow under the von Neumann Machine
  - Reduce data access time under current systems
  - It is a new system paradigm based on performance modeling
- **Deliverable:** Scalable data management systems
  - The Hermes system for *Dataflow<sub>V</sub>*
  - Performance engineering & performance-guided system design
- **Keywords:**
  - Data centric rethinking, Data system optimization, HPC4AI
- **Challenges:**
  - Understanding and Support of data concurrency and system integration (the ecosystem)
  - Hardware/software co-design
  - **Infrastructure and testbed**

**Applications & Infrastructure**



# Conclusion (scalable data management)

---

- A new computer paradigm is introduced for AI-DL & big data
  - **Dataflow under the von Neumann Machine**
- Performance modeling and understanding is essential for the success of *Dataflow<sub>V</sub>*
- Current Success: the Hermes System
  - Concurrent data access and system integration
- Ongoing Project: IOWrap
  - Geared up for AI applications
  - StoreHub: A community and community testbed

SOFTWARE & HARDWARE CODESIGN FOR  
DATAFLOW UNDER VON NEUMANN MACHINE



# Key members of current research staff & students at Illinois Tech

**Dr. Kun Feng**  
Research Software Engineer

**Eneko Gonzalez**  
Research Software Engineer

**Xiaoyang Lu**  
Research Professor

The **HDF Group**



**Anthony Kougkas,**  
Research Professor



**Sandia  
National  
Laboratories**



**Keith Bateman**  
PhD Candidate

**Vadim Biryukov**  
PhD Student

**Jaime Cernuda**  
PhD Candidate

**Luke Logan**  
PhD Candidate

**Isa Muradli**  
PhD Student



**Neeraj Rajesh**  
PhD Candidate

**Meng Tang**  
PhD Candidate

**Hua Xu**  
PhD Student

**Jie Ye**  
PhD Candidate

**Izzet Yildirim**  
PhD Candidate



***Thank you  
Any questions?***



## AI & Data: Challenge & Opportunity in Computer System Research

Xian-He Sun

The Gnosis Research Center at the Illinois Institute of Technology

[sun@iit.edu](mailto:sun@iit.edu)  
[www.cs.iit.edu/~scs](http://www.cs.iit.edu/~scs)



We would like to thank  
our sponsors the  
National Science  
Foundation

